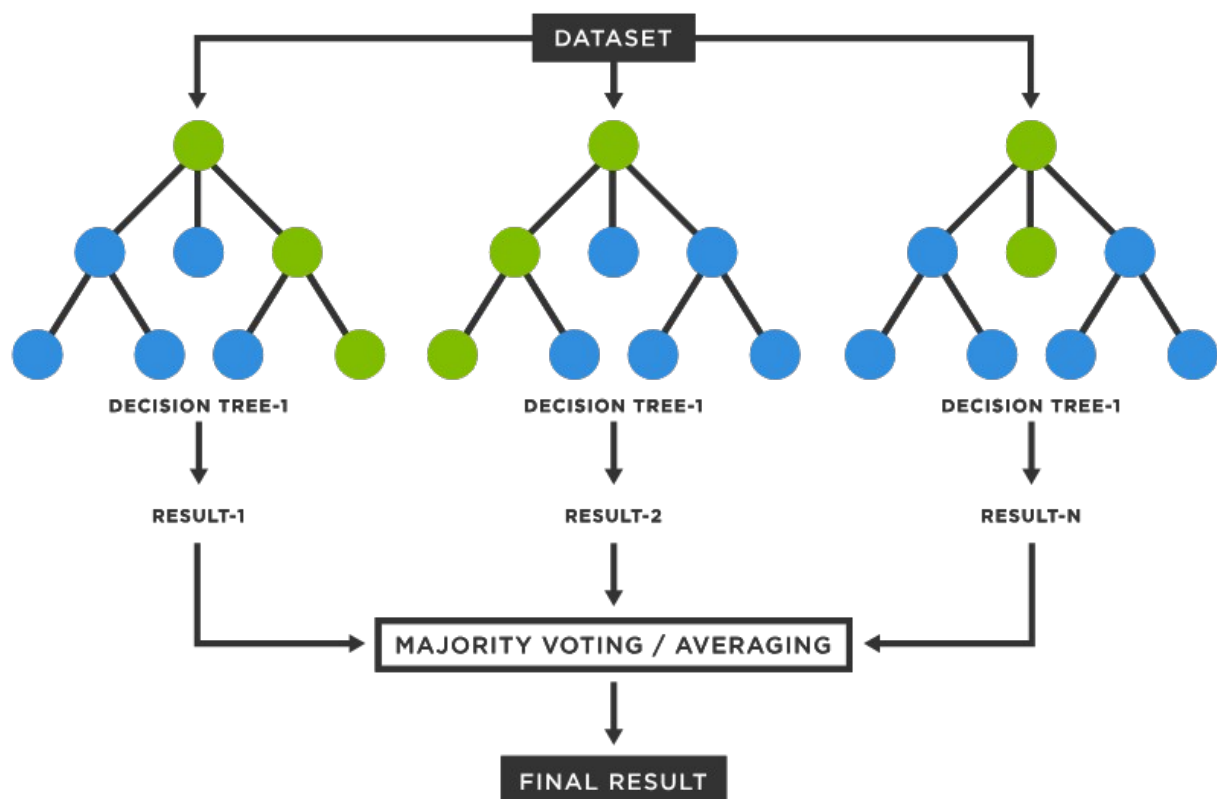


Customer churn refers to the situation where a customer discontinues their association with a particular company. There can be various reasons behind this decision, such as affordability concerns, dissatisfaction with the product or service, or inadequate customer support.

Frequently, customers who churn from one company tend to move their business to a competitor. For instance, if you are unsatisfied with your current mobile service provider due to slow Internet speed, you are more likely to switch to an alternative provider.

The process of churning typically doesn't happen abruptly. If you encounter issues like low network bandwidth, you may endure it for a month or two. During this period, you might contact customer support, assess your network speed, and even share your discontent on social media through reviews.

- The dataset used in this project: <https://www.kaggle.com/datasets/blastchar/telco-customer-churn>
- Dependencies: pandas, matplotlib, seaborn, scikit-learn, Imblearn, colorama | Use the "pip install" command to install the modules
- The ML Model used in this project is the Random Forests model: is a supervised machine learning algorithm. It is one of the most used algorithms due to its accuracy, simplicity, and flexibility. The fact that it can be used for classification and regression tasks, combined with its nonlinear nature, makes it highly adaptable to a range of data and situations.



ew.

```
import pandas as pd

# reading data
df = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn.csv")

# Show the data
df.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure
0	7590-VHVEG	Female	0	Yes	No	1
1	5575-GNVDE	Male	0	No	No	34
2	3668-QPYBK	Male	0	No	No	2
3	7795-CF0CW	Male	0	No	No	45
4	9237-HQITU	Female	0	No	No	2

	MultipleLines	InternetService	OnlineSecurity	...
DeviceProtection \				
0	No phone service	DSL	No	...
No				
1	No	DSL	Yes	...
Yes				
2	No	DSL	Yes	...
No				
3	No phone service	DSL	Yes	...
Yes				
4	No	Fiber optic	No	...
No				

	TechSupport	StreamingTV	StreamingMovies	Contract
PaperlessBilling \				
0	No	No	No	Month-to-month
Yes				
1	No	No	No	One year
No				
2	No	No	No	Month-to-month
Yes				
3	Yes	No	No	One year
No				
4	No	No	No	Month-to-month
Yes				

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

[5 rows x 21 columns]

```
import colorama as color
```

```
green_color = color.Back.GREEN
```

```
red_color = color.Back.RED
```

```
reset_color_fore = color.Fore.RESET
```

```
# Shape of the data
```

```
shape = df.shape
```

```
print(green_color + reset_color_fore + "Number of rows: " +  
str(shape[0]), "Number of Columns: " + str(shape[1]))
```

```
Number of rows: 7043 Number of Columns: 21
```

```
# Information about the data
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines           7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup            7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies         7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling        7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                  7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

○ Each user is identified through a unique customer ID. There are 19 independent variables used to predict the target feature – customer churn. In this dataset, customer churn is defined as users who have left within the last month.

```

# We can see that there are no missing values
if df.isnull().sum().sum() == 0:
    print(red_color + "There are no missing values...")

There are no missing values...

# Number of customers who have churned
number_of_customer_churned = df["Churn"].value_counts()
number_of_customer_churned

Churn
No      5174
Yes     1869
Name: count, dtype: int64

not_churned = (number_of_customer_churned[0] / shape[0] * 100)
churned = "%.0f" % (100 - not_churned)
print("%.0f" % not_churned + f"% of customers who have not churn and
{churned}% of customers who have churned.")

```

73% of customers who have not churn and 27% of customers who have churned.

☉ We are dealing with an imbalanced classification problem. We will need to perform some feature engineering to create a balanced training dataset before building the predictive model.

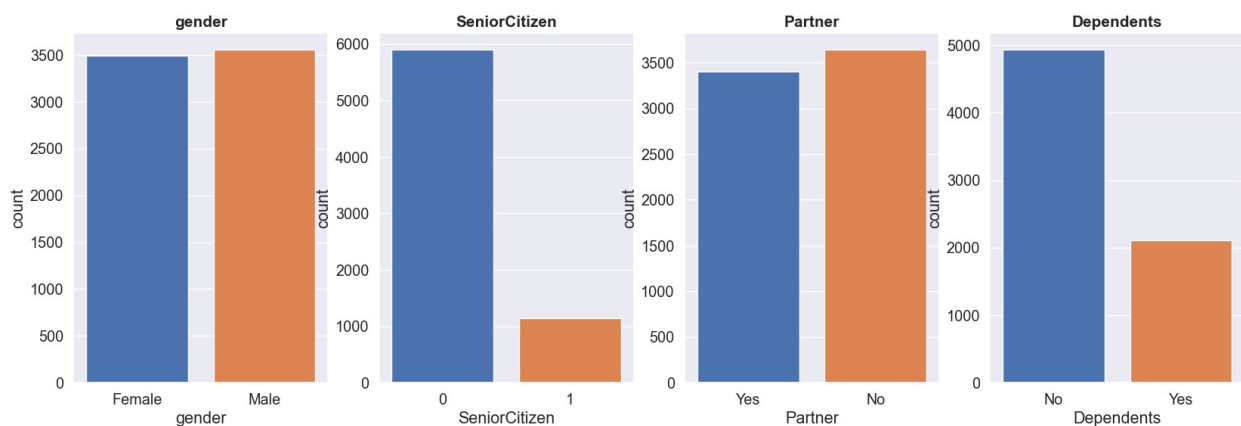
Exploratory Data Analysis for Customer Churn Prediction

```
import matplotlib.pyplot as plt
import seaborn as sea
import numpy as np

# Analyzing demographic customer data
demographic_cols = ['gender', 'SeniorCitizen', 'Partner',
                    "Dependents"]

# Change the figure size
plt.figure(figsize=(20, 6))

sea.set(font_scale = 1.3)
for idx, col in enumerate(demographic_cols):
    ax = plt.subplot(1, len(demographic_cols), idx + 1)
    sea.countplot(x=col, data=df, saturation=1)
    ax.set_title(str(col), fontdict={'size': 15, 'weight': 'bold'})
```



```
# Change the figure size
plt.figure(figsize=(8, 5))
sea.set(font_scale = 0.9)

# Get the current axes object
ax = plt.gca()

# Set the font weight of x and y tick labels
for xtick_label in ax.get_xticklabels():
    xtick_label.set_weight("bold")
```


```

for ytick_label in ax.get_yticklabels():
    ytick_label.set_weight("bold")

# Create a boxplot
sea.boxplot(x='Churn', y='MonthlyCharges', data=df,
palette=['#a343a2', "#FAAE7B"], linewidth=3)

<Axes: xlabel='Churn', ylabel='MonthlyCharges'>

```



○ we can see that there are no outliers in the data

```

# analyze the relationship between customer churn and a few other
categorical variables captured in the dataset

cols = ['InternetService', "TechSupport", "OnlineBackup", "Contract"]

plt.figure(figsize=(20, 5))

for i, col in enumerate(cols):
    ax = plt.subplot(1, len(cols), i+1)
    sea.countplot(x="Churn", hue = str(col), data = df)
    ax.set_title(f"{col}", fontdict={'size': 15, 'weight': 'bold'})


# Get the current axes object

```

```
ax = plt.gca()

# Set the font-weight of x and y tick labels
for xtick_label in ax.get_xticklabels():
    xtick_label.set_weight("bold")

for ytick_label in ax.get_yticklabels():
    ytick_label.set_weight("bold")
```



Let's look into each attribute:

- **InternetService:** It is clear from the visual above that customers who use fiber optic Internet churn more often than other users. This might be because fiber Internet is a more expensive service, or this provider doesn't have good coverage.
- **TechSupport:** Many users who churned did not sign up for tech support. This might mean that these customers did not receive any guidance on fixing technical issues and decided to stop using the service.
- **OnlineBackup:** Many customers who had churned did not sign up for an online backup service for data storage.
- **Contract:** Users who churned were almost always on a monthly contract. This makes sense, since these customers pay for the service on a monthly basis and can easily cancel their subscription before the next payment cycle.

Even without building a fancy machine learning model, a simple data-driven analysis like this can help organizations understand why they are losing customers and what they can do about it.

For instance, if the company realizes that most of their users who churn have not signed up for tech support, they can include this as a complimentary service in some of their future product offerings to prevent other customers from leaving.

Preprocessing Data for Customer Churn

```
# Convert the the TotalCharges column from object to numeric column
df['TotalCharges'] = df['TotalCharges'].apply(lambda x:
pd.to_numeric(x, errors='coerce')).dropna()
```

Encoding the categorical data

```
cat_features =
df.drop(['customerID', 'TotalCharges', 'MonthlyCharges', 'SeniorCitizen',
'tenure'], axis=1)
cat_features.head()
```

	gender	Partner	Dependents	PhoneService	MultipleLines
0	Female	Yes	No	No	No phone service
1	Male	No	No	Yes	No
2	Male	No	No	Yes	No
3	Male	No	No	No	No phone service
4	Female	No	No	Yes	No

	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV
0	No	Yes	No	No	No
1	Yes	No	Yes	No	No
2	Yes	Yes	No	No	No
3	Yes	No	Yes	Yes	No
4	No	No	No	No	No

	StreamingMovies	Contract	PaperlessBilling
0	No	Month-to-month	Yes
1	No	One year	No
2	No	Month-to-month	Yes
3	No	One year	No
4	No	Month-to-month	Yes

	Churn
0	No
1	No
2	Yes


```
3 No
4 Yes
```

```
from sklearn import preprocessing
```

```
le = preprocessing.LabelEncoder()
df_cat = cat_features.apply(le.fit_transform)
df_cat.head()
```

	gender	Partner	Dependents	PhoneService	MultipleLines
InternetService \					
0	0	1	0	0	1
0					
1	1	0	0	1	0
0					
2	1	0	0	1	0
0					
3	1	0	0	0	1
0					
4	0	0	0	1	0
1					

	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport
StreamingTV \				
0	0	2	0	0
0				
1	2	0	2	0
0				
2	2	2	0	0
0				
3	2	0	2	2
0				
4	0	0	0	0
0				

	StreamingMovies	Contract	PaperlessBilling	PaymentMethod	Churn
0	0	0	1	2	0
1	0	1	0	3	0
2	0	0	1	3	1
3	0	1	0	0	0
4	0	0	1	2	1

```
num_features =
df[['customerID', 'TotalCharges', 'MonthlyCharges', 'SeniorCitizen', 'tenure']]
finaldf = pd.merge(num_features, df_cat, left_index=True,
right_index=True)
finaldf.head(3)
```

	customerID	TotalCharges	MonthlyCharges	SeniorCitizen	tenure
gender \					
0	7590-VHVEG	29.85	29.85	0	1
0					
1	5575-GNVDE	1889.50	56.95	0	34
1					
2	3668-QPYBK	108.15	53.85	0	2
1					

	Partner	Dependents	PhoneService	MultipleLines	...
OnlineSecurity \					
0	1	0	0	1	...
0					
1	0	0	1	0	...
2					
2	0	0	1	0	...
2					

	OnlineBackup	DeviceProtection	TechSupport	StreamingTV
StreamingMovies \				
0	2	0	0	0
0				
1	0	2	0	0
0				
2	2	0	0	0
0				

	Contract	PaperlessBilling	PaymentMethod	Churn
0	0	1	2	0
1	1	0	3	0
2	0	1	3	1

[3 rows x 21 columns]

we will use a technique called oversampling. This is a process that involves randomly selecting samples from the minority class and adding it to the training dataset. We are going to oversample the minority class until the number of data points are equal to that of the majority class.

```
from sklearn.model_selection import train_test_split

finaldf = finaldf.dropna()
finaldf = finaldf.drop(['customerID'],axis=1)

X = finaldf.drop(['Churn'],axis=1)
y = finaldf['Churn']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.33, random_state=42)
```

```

from imblearn.over_sampling import SMOTE

# Class to perform over-sampling using SMOTE.
oversample = SMOTE(k_neighbors=5)
X_smote, y_smote = oversample.fit_resample(X_train, y_train)
X_train, y_train = X_smote, y_smote

# Check the number of samples in each class to ensure that they are equal
# There should be 3,452 values in each class, which means that the training dataset is now balanced.
y_train.value_counts()

Churn
0      3452
1      3452
Name: churn, dtype: int64

# Creating and fitting the model
from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(random_state=46)
rf.fit(X_train, y_train)

RandomForestClassifier(random_state=46)

# Testing the accuracy of the model
from sklearn.metrics import accuracy_score

preds = rf.predict(X_test)
print(accuracy_score(preds, y_test))

0.7703576044808272

```

The model accuracy is around 0.77%. It is not so high accuracy but for the dataset used in this project it is a good accuracy